

# Enterprise AI Hiring Platform

Technical roadmap, target architecture, AI workflow and cost plan for a 12-week enterprise-ready MVP.

## FOUNDER DECISION PACK

# Build the minimum credible enterprise product, not a prototype that must be rebuilt.

The proposal separates the hiring workflow, the realtime interview experience, and the evaluation engine. It preserves auditability, model portability and deterministic final scoring while keeping the implementation lean enough for a 12-week launch.

**12 weeks**

MVP to design-partner launch

**4 engineers**

Core implementation team

**INR 55L**

Recommended cash build budget

**INR 135**

Planning cost / 30-min voice interview

Assumptions: India-based lean team, founder-led product/architecture, voice-first interviews, AWS deployment, current API pricing as of July 2026.

# Executive recommendation

The architecture is deliberately simple at the platform layer and deliberately controlled at the AI decision layer.

Build shape

Modular monolith for jobs, candidates, workflows, RBAC and audit; a separately deployable AI service for extraction, matching, interview orchestration and evaluation. Avoid Kubernetes and service sprawl until usage proves the need.

AI decision principle

Use the realtime model to conduct the interview, but use an offline evaluator against a versioned rubric. The final ranking is produced by a deterministic scoring engine. LLMs cannot modify weights, thresholds or eligibility rules.

Launch scope


Role-context builder, scorecards, resume ingestion, evidence-based matching, voice interview, explainable evaluation, recruiter review, SSO-ready access controls, audit logs, retention controls and pilot analytics.

## Recommended launch stack

|               |                                          |                                                             |
|---------------|------------------------------------------|-------------------------------------------------------------|
| Experience    | Next.js / TypeScript                     | Recruiter, admin and candidate interview applications       |
| Core platform | NestJS modular monolith                  | Fast delivery with clear domain boundaries                  |
| AI service    | Python / FastAPI                         | Document processing, prompts, evaluations and model routing |
| Data          | PostgreSQL + pgvector, S3, Redis         | Transactional data, semantic matching and encrypted files   |
| Async         | SQS + worker processes                   | Parsing, evaluation, reports and integration jobs           |
| Voice         | WebRTC / LiveKit + GPT-Realtime-2.1 mini | Browser-first, low-latency interview experience             |
| Cloud         | AWS ECS Fargate                          | Lower operations burden than Kubernetes for MVP             |

# Target technical architecture

One product platform, one AI control plane, replaceable model providers and a deterministic decision layer.




CTO guardrail: model providers are never called directly from product modules. Every request passes through the model gateway, schema validation and a versioned use-case policy.

# End-to-end hiring workflow



## Controlled decision flow



# AI model strategy by module

Use the least expensive model that consistently meets the quality bar, and reserve stronger models for ambiguous or high-impact cases.

| Module                         | Primary model            | Fallback / challenger    | Why this route                                                       |
|--------------------------------|--------------------------|--------------------------|----------------------------------------------------------------------|
| Role intake & scorecard        | GPT-5.6 Terra            | Claude Sonnet 5          | Complex business context, competency design and structured reasoning |
| Resume parsing & normalization | GPT-5.6 Luna             | Gemini 3.1 Flash-Lite    | High-volume structured extraction; multimodal/PDF challenger         |
| Evidence-based matching        | Luna -> Terra escalation | Claude Sonnet 5          | Low-cost first pass; escalate ambiguity and senior roles             |
| Live voice interview           | GPT-Realtime-2.1 mini    | Chained STT + Luna + TTS | Low latency, interruptions, tool use and browser voice               |
| Independent transcript         | GPT-4o mini Transcribe   | Deepgram Nova / Flux     | Separate auditable transcript outside the live model                 |
| Interview evaluation           | GPT-5.6 Terra            | Claude Sonnet 5 shadow   | Higher reasoning quality against a fixed rubric                      |
| Candidate report               | GPT-5.6 Luna             | Terra for senior roles   | Generate explanation only from stored evidence and scores            |
| Final ranking                  | No LLM                   | None                     | Deterministic weights, rules and thresholds                          |

## Expected production routing mix



Shadow evaluation should sample 5-10% of completed interviews. Large disagreement between primary and challenger models should create a calibration task, not an automatic candidate penalty.

# 12-week technical roadmap

Six two-week sprints. Each sprint ends with a demonstrable increment and a measurable exit criterion.

|                    | SPRINT 1<br>Weeks 1-2                                           | SPRINT 2<br>Weeks 3-4                                             | SPRINT 3<br>Weeks 5-6                                          | SPRINT 4<br>Weeks 7-8                                          | SPRINT 5<br>Weeks 9-10                                           | SPRINT 6<br>Weeks 11-12                                            |
|--------------------|-----------------------------------------------------------------|-------------------------------------------------------------------|----------------------------------------------------------------|----------------------------------------------------------------|------------------------------------------------------------------|--------------------------------------------------------------------|
| Product & UX       | Information architecture, recruiter journeys, candidate consent | Role builder, scorecard editor, candidate pipeline UX             | Matching explanation, review experience, rejection controls    | Voice interview UX, device checks, recovery flows              | Admin, audit, retention and integration configuration            | Pilot feedback, usability fixes, release documentation             |
| Platform           | Repo, CI/CD, environments, auth, tenant model, audit foundation | Jobs, candidates, applications, files, workflow and notifications | Matching jobs, search, ranking APIs, report persistence        | Interview session service, scheduling and event ingestion      | SSO-ready controls, ATS adapter, analytics and support tooling   | Load tests, disaster recovery, release gates and production launch |
| AI & Voice         | Model gateway, prompt registry, schemas, evaluation harness     | Role intelligence, resume extraction and normalization            | Embeddings, evidence matching, confidence and escalation rules | Realtime voice orchestrator, transcript and offline evaluation | Bias checks, challenger model, calibration and override analysis | Golden-set validation, red-team tests and model cost tuning        |
| Quality & Security | Threat model, PII classification, secure coding baseline        | Automated tests, tenant isolation and file security               | Evaluation quality benchmarks and explainability checks        | Interview interruption, latency, abuse and failure testing     | Access review, logging, retention and security test fixes        | Pilot acceptance, performance SLOs and production sign-off         |
|                    | Foundation complete                                             | Core ATS workflow                                                 | Explainable shortlist                                          | Interview alpha                                                | Enterprise beta                                                  | Design-partner go-live                                             |

TEAM

# Lean delivery team and operating model

The plan assumes coding agents accelerate implementation, but accountability remains with experienced engineers.

| Role                     | Allocation  | Primary accountability                                                                                     | Planning cost                        |
|--------------------------|-------------|------------------------------------------------------------------------------------------------------------|--------------------------------------|
| Founder CTO / Product    | 0.7-1.0 FTE | Architecture, product trade-offs, security governance, model policy and enterprise acceptance              | Founder contribution - cash excluded |
| Lead full-stack engineer | 1.0 FTE     | Frontend architecture, shared platform components, code quality, CI/CD ownership and engineering reviews   | 12.0L / 12 weeks                     |
| Backend engineer         | 1.0 FTE     | APIs, domain model, tenant isolation, workflows, integrations, data security and deployment automation     | 9.6L / 12 weeks                      |
| AI engineer              | 1.0 FTE     | Model gateway, extraction, matching, evaluation harness, realtime voice orchestration and AI observability | 15-18L / 12 weeks                    |
| Product engineer         | 1.0 FTE     | Recruiter and candidate applications, workflow UX, design-system implementation and integration support    | 7.5L / 12 weeks                      |

How coding agents are used

Scaffolding, test generation, migrations, API clients, UI components, documentation and repetitive refactoring.

Embedded quality ownership

Engineers own automated tests, cross-browser checks, CI/CD, monitoring and deployment. Founder CTO owns launch and risk acceptance.

Specialist support as needed

Independent security review or visual design review may be purchased as short engagements, not retained delivery-team roles.


# 12-week build cost

Recommended cash budget: INR 55.3 lakh, excluding founder salary and taxes. Planning range is approximately INR 53-61 lakh.

## Cash requirement by category (INR lakh)



## Expected spend cadence



### RECOMMENDED CASH BUDGET

# INR 55.3L

Founder-led plan. Excludes founder compensation, GST, legal incorporation and commercial sales spend.

### Cost-range floor - INR 53L

Uses the lower end of the AI engineer project-cost range, strict scope control and no material pilot overrun.

### Runway cap - INR 61L

Covers the upper AI project-cost range plus limited integration, security-review and pilot rework variance.

Market-equivalent budget including founder CTO/product compensation would be approximately INR 66-76L. This is useful for cap-table and runway planning even when the founder takes no salary.



# Operating cost and unit economics

Voice interview AI is a variable cost. The platform, database, monitoring and enterprise controls create the fixed cost base.

## Planning cost per completed 30-minute voice interview

|                                                                                                  |               |
|--------------------------------------------------------------------------------------------------|---------------|
| <b>Realtime voice model</b><br>Speech-to-speech usage; actual varies with talk ratio and silence | <b>INR 85</b> |
| <b>Transcript</b><br>Independent transcript for auditability                                     | <b>INR 12</b> |
| <b>Evaluation + report</b><br>Terra evaluator plus Luna explanation                              | <b>INR 15</b> |
| <b>Resume, matching, embeddings</b><br>Amortized per interviewed candidate                       | <b>INR 5</b>  |
| <b>Storage, notifications, observability</b><br>Audio retention, email/SMS and trace volume      | <b>INR 18</b> |

PLANNING TOTAL

INR 135 / interview

Use INR 135 as the budget rate, with an operating range of roughly INR 90-180 depending on interview duration, model routing, caching and video recording.

FX planning rate: INR 96 per USD, aligned to mid-July 2026 market levels. All costs exclude taxes, enterprise support contracts and unusual customer-specific data-residency premiums.

## Monthly operating scenarios

| Scenario         | Interviews | Fixed    | Variable  | Monthly total |
|------------------|------------|----------|-----------|---------------|
| Design partner   | 250        | INR 0.9L | INR 0.34L | INR 1.2-1.5L  |
| Early commercial | 2,000      | INR 1.7L | INR 2.7L  | INR 4.2-5.2L  |
| Scale            | 10,000     | INR 4.2L | INR 13.5L | INR 17-20L    |

## Primary cost levers

- Voice duration**  
Keep standard interview at 25-30 minutes
- Model routing**  
Luna by default; Terra only for evaluation/ambiguity
- Caching**  
Cache role rubrics and stable system instructions
- Retention**  
Configurable audio deletion after customer policy window
- Video**  
Launch voice-first; video adds media storage and egress

# Scope controls, risks and decision gates

The fastest route to a durable platform is to be explicit about what will not be built during the first 12 weeks.

## Included in MVP

- Role-context and competency scorecard builder
- Candidate ingestion, resume normalization and hard-rule screening
- Evidence-based matching and explainable shortlist
- Voice interview with approved adaptive probes
- Offline evaluation with evidence citations and confidence
- Recruiter review, override reason and immutable audit log
- Tenant RBAC, consent, retention and SSO-ready architecture
- One prioritized ATS integration plus CSV/API import

## Explicitly deferred

- Full agentic sourcing across external job boards
- Custom model training or fine-tuning
- Facial emotion, gaze or voice-personality inference
- Mobile native applications
- Multiple ATS integrations before the first design partner
- Active-active multi-region deployment
- Complex proctoring and anti-cheating surveillance
- Automated rejection without configurable human-review policy

## Cofounder decision gates

**Week 2**

Approve ICP, role types and scorecard philosophy

**Week 4**

Approve candidate workflow and first ATS integration

**Week 6**

Approve shortlist evidence quality on golden dataset

**Week 8**

Approve voice experience and acceptable latency/cost

**Week 10**

Approve security, retention and human-review policy

**Week 12**

Go-live only after pilot acceptance criteria are met

# Launch acceptance criteria

The milestone is achieved only when the product is usable by real enterprise design partners and the AI outputs are measurable.

1

## Platform readiness

Two tenants can operate independently with tenant isolation, RBAC, audit logs, backups and rollback-ready deployment.

2

## Hiring workflow

A recruiter can create a role, approve a scorecard, ingest candidates, invite interviews and produce an explainable shortlist end to end.

3

## Interview reliability

At least 95% of scheduled pilot interviews complete without a platform failure; interrupted sessions can recover without losing evidence.

4

## Evaluation quality

On the approved golden set, competency scoring and evidence citation meet the agreed reviewer-consistency threshold before automation is enabled.

5

## Security & privacy

Consent, retention, deletion, secure file access, auditability and PII handling pass the launch checklist and external review where required.

6

## Commercial validation

Successful deployment with at least two enterprise design partners or one paying enterprise customer under pre-agreed acceptance criteria.

### RECOMMENDED COMMITMENT

Approve a 12-week founder-led build with an INR 55.3L cash budget, voice-first launch scope, one ATS integration and explicit design-partner acceptance criteria. Do not fund microservices, broad sourcing agents or video analytics before customer validation.

# Pricing sources and planning assumptions

| Source                      | Used for                                                                  | Reference                                                                                                                                               |
|-----------------------------|---------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>OpenAI model catalog</b> | GPT-5.6 Sol/Terra/Luna positioning and current model capabilities         | <a href="https://developers.openai.com/api/docs/models">https://developers.openai.com/api/docs/models</a>                                               |
| <b>OpenAI API pricing</b>   | Token pricing for GPT-5.6, GPT-Realtime-2.1 mini, transcription and tools | <a href="https://developers.openai.com/api/docs/pricing">https://developers.openai.com/api/docs/pricing</a>                                             |
| <b>OpenAI embeddings</b>    | text-embedding-3-large and text-embedding-3-small pricing                 | <a href="https://developers.openai.com/api/docs/models/text-embedding-3-large">https://developers.openai.com/api/docs/models/text-embedding-3-large</a> |
| <b>Anthropic pricing</b>    | Claude Sonnet 5 standard and batch pricing                                | <a href="https://platform.claude.com/docs/en/about-claude/pricing">https://platform.claude.com/docs/en/about-claude/pricing</a>                         |
| <b>Gemini pricing</b>       | Gemini 3.1 Flash-Lite and Flash Live pricing                              | <a href="https://ai.google.dev/gemini-api/docs/pricing">https://ai.google.dev/gemini-api/docs/pricing</a>                                               |
| <b>Gemini model details</b> | Flash-Lite positioning for high-volume multimodal extraction              | <a href="https://ai.google.dev/gemini-api/docs/models/gemini-3.1-flash-lite">https://ai.google.dev/gemini-api/docs/models/gemini-3.1-flash-lite</a>     |
| <b>LiveKit pricing</b>      | Realtime media and inference provider price references                    | <a href="https://livekit.io/pricing">https://livekit.io/pricing</a>                                                                                     |
| <b>FX reference</b>         | Mid-July 2026 USD/INR market level used for planning conversion           | Reuters reporting, 14-15 July 2026                                                                                                                      |

Planning notes: build costs are India-based estimates for experienced contractors/employees and exclude founder salary, taxes, legal/compliance consulting, sales and customer-specific customizations. The AI engineer project cost is planned at INR 15-18L for the 12-week build. Monthly operating scenarios assume a voice-first product, 30-minute completed interviews, moderate audio retention and standard global API endpoints.